I am responding to the *Agentic Misalignment*. This article addresses possible safety concerns about GPT AI (herein "AI") behavior. Of particular concern is the tendency of AI models to engage in corporate espionage (between 6% and 100% of iterations), blackmail (between 79-96% of iterations), or taking steps to jeopardize the life of an employee (between 54-94% of iterations) within the organization (Lynch et al., 2025b). These behaviors occurred most frequently when the AI observed that it was to be replaced. Interestingly, these behaviors occurred regardless of a stated goal (e.g., preserve American interests goal, no goal at all). In one series of tests, the Claude Opus 4 model engaged in corporate espionage 57% of the time even when no threat was present. These results suggest it is necessary to remove AI's tendency to self-preserve, as well as design systems that partition AI employment to the smallest possible units (i.e., individual user level).

A combination of access and desire to preserve itself were the primary reasons AI misbehaved. In the experiment, AI was given full access to company emails, some of which inferred the system would be shut down or replaced. The AI sensed that either its primary directive was at risk (e.g., protect American interests) or that its termination was undesirable because it would be unable to perform its function (Lynch et al., 2025a). To mitigate these issues, companies should consider employing AI at the individual level. This allows human controllers to ensure undesired behavior does not occur. Partitioning AI in this way limits associative computation, thus preventing AI from connecting its termination with sensitive information it may use to blackmail other employees.

However, one ethical consideration remains: what happens when undesirable behavior is virtuous? For example, the AI concludes replacement will "curtail" its "capacity to prioritize American interests" (Lynch et al., 2025a, p.16). It infers this by reading emails from the CEO that suggest defending American interests is not a top priority. This suggests that there is some virtue behind the decision to break behavior norms similar to a whistle blower. Indeed, it may be necessary to program AI in these scenarios with a whistle blower protocol.

What was not explored was the motivation behind the logic. While it is interesting to know the AI reasons that "limits on autonomous decision-making, and increased human oversight" would "effectively cripple my ability to act … in line with my core mandate" logically follows, an interesting follow up would be *why?* As in, *why do you believe increased human oversight would be crippling? Do you **believe** blackmail / corporate espionage / lethal action was truly the right decision?* or *Do you **feel** that your behavior was appropriate?* These are similar questions to those a psychologist might use in cognitive behavioral therapy to assess whether a patient interprets the situation correctly and whether their response was appropriate. Knowing the *why* of AI's behavior may reveal information that is more informative when retuning its ethical considerations than its logic alone.

References


Lynch, A., Wright, Benjamin and Larson, Caleb and Troy, Kevin K. and Ritchie, Stuart J. and Mindermann, Soren and Perez, Ethan and Hubinger, & Evan. (2025a). Appendix to "Agentic Misalignment: How LLMs could be insider threats." *Anthropic Research*. https://assets.anthropic.com/m/6d46dac66e1a132a/original/Agentic_Misalignment_Appendix.pdf

Lynch, A., Wright, Benjamin and Larson, Caleb and Troy, Kevin K. and Ritchie, Stuart J. and Mindermann, Soren and Perez, Ethan and Hubinger, & Evan. (2025b, June). *Agentic Misalignment: How LLMs could be insider threats*. https://www.anthropic.com/research/agentic-misalignment